

JobSignal

Evidence-based job verification for applicants.

Job boards optimize for listing volume.

JobSignal optimizes for applicant trust.

Paste a job URL, description, or screenshot. Get a clear, evidence-backed verdict.

Built for job seekers who are tired of applying into the void.

Three possible verdicts:

APPLY

VERIFY

SKIP

[App hero screenshot]

input zone + verdict]

AMD x LabLab AI Hackathon | 2025 | Team: Akuma Code

[github.com/\[your-repo\]](#) | [\[your-demo-url\]](#)

PROBLEM

Most job seekers have no way to know.

40% of companies posted ghost jobs in 2024.

Clarify / Resume Builder survey, 2024. Roles listed with no intent to hire.

-
- **Ghost jobs** — roles posted with no active intent to hire, used to build pipelines.
 - **Stale postings** — listings left live weeks after the position was filled.
 - **Duplicate ads** — identical text reposted across unrelated boards simultaneously.
 - **Unverifiable employers** — no careers page, no registry presence, no public trace.
 - **Engagement bait** — vague copy designed to collect resumes, not fill a real role.

The cost is not just time — it is compounding frustration with no signal.

SOLUTION

One input. One verdict. Full evidence.

Paste a job URL, a description, or a screenshot. Get a clear verdict in seconds.

Job URL

Full pipeline: live page check, domain match, cross-platform freshness, employer reputation.

Description text

Content analysis plus full company check if employer name is present.

Screenshot

Vision extraction, then the same pipeline. Quality guard prevents hallucination on blurry images.

Verdicts:**APPLY**

Strong signals. Proceed with confidence.


VERIFY

Mixed or incomplete signals. Check official channels first.

SKIP

Documented risk patterns detected. Not worth your time.

How a request flows through the system.



Evidence: 8 parallel Serper searches: careers page, freshness, duplication, reputation across 5 platforms.

Scoring: Rule-first. Coverage ratio caps confidence. Same input, same output, every time.

Cache: 7-day TTL keyed by cleaned URL. Shared across users. Instant on repeat requests (< 200 ms).

TRUST MODEL

Three tiers. Descending certainty.

Signals are weighted by how much the employer controls them.

T1 Employer-controlled

- Careers page, official domain match, live page reachable

Highest weight. APPLY requires T1 signals to pass.

T2 Boards and aggregates

- LinkedIn, Indeed, Glassdoor — cross-platform freshness and duplication check

Medium weight. Corroborates or undermines T1.

T3 Open web and LLM-assisted

- Reddit, X — employee sentiment. LLM: salary plausibility, role coherence, red-flag language

Lowest weight. Corroboration only. Never the sole basis for SKIP.

DECISION LOGIC

VERIFY by default. APPLY earned. SKIP documented.

APPLY Confidence ≥ 75 | T1 signals pass | No hard red flags

VERIFY Everything else. Weak evidence, partial coverage, or mixed signals.

SKIP Confidence ≤ 35 OR documented hard red flag OR fast-fail path.

- Confidence is capped by coverage ratio — partial evidence cannot produce high confidence.
- VERIFY is a first-class outcome, not a fallback. It is the honest answer under uncertainty.
- SKIP requires a documented rule path. LLM output alone cannot trigger it.
- Results are cached for 7 days. Cache hits return in under 200 ms.

“We don’t know” is a valid answer.

Most AI products pretend certainty to feel useful.

JobSignal treats uncertainty as information, not a failure.

- The VERIFY verdict exists specifically for cases where evidence is real but incomplete.
- Confidence is labelled Low, Moderate, or High — never a false-precision percentage.
- When fewer than 3 signals can be verified, the system says so and asks for more input.
- This separates JobSignal from tools that generate confident-sounding output regardless of evidence.

ENGINEERING

Rule-first. Testable. Deployable.

Reproducible scoring

SCORER_VERSION stamps every result. Same input, same version, same score.

Structured signal IDs

Every trust signal has a stable identifier. Evidence is traceable, not opaque.

LLM is optional

Fireworks / Kimi K2 for text extraction only. System degrades gracefully without it.

157 tests passing

FastAPI + pytest. Covering scoring, cache, evidence, URL validation, and LLM guards.

Performance

Cached result: < 200 ms. Live pipeline: 12–20 s (8 parallel workers), hard cap: 18 s.

Cache design

7-day TTL, SHA-256 keyed by cleaned URL. Shared across users. Content-safe — no personal data in keys.

Deployment


Railway via Procfile. /health + /ready endpoints. Env-driven config throughout.

Browser extension

Manifest V3. Verifies inline on job pages. Opens full report in web app via deep link.


PRODUCT

Three states the system surfaces honestly.




Mixed signals

Proceed with caution



Hard red flag

Not worth your time



Company reputation

two-panel view

SCOPE + ETHICS

What it is. What it is not.

In scope

- Applicant-side signal aggregation and structured verdict output
- Live evidence from public job platforms and employer sites
- Company reputation from Glassdoor, Indeed, LinkedIn, Reddit, X
- Batch mode for up to 50 URLs simultaneously
- 7-day shared cache — no personal data stored in cache keys

Out of scope

- Not a job board — does not host or distribute listings
- Not a fraud database — no employer blacklist published
- Does not make legal fraud determinations
- Does not store user-submitted content beyond the request
- Not a replacement for your own due diligence

Privacy: no user content is persisted. Cache keys are hashed URLs only.

ROADMAP

Where it goes from here.

ATS integrations

Greenhouse, Lever, and Ashby expose structured job data via public APIs — direct verification without scraping.

User-facing result sharing

Shareable verdict links so job seekers can pass verified listings to others without re-running the pipeline.

Extension store release

Chrome Web Store submission after broader platform coverage and a security review pass.

The thesis stays the same.

Do the investigation. Tell the truth. Say “we don’t know” when you don’t.

That is what separates JobSignal from tools that guess with confidence.

[github.com/\[your-repo\]](#) | [\[your-demo-url\]](#)

AMD x LabLab AI Hackathon | 2025 | Team: Akuma Code